



# Java Micro-Benchmarking

Constantine Nosovsky



# Agenda

- Benchmark definition, types, common problems
- Tools needed to measure performance
- Code warm-up, what happens before the steady-state
- Using JMH
- Side effects that can affect performance
- JVM optimizations (Good or Evil?)
- A word about concurrency
- Full example, “human factor” included
- A fleeting glimpse on the JMH details

# Benchmark definition, types, common problems





# What is a “Benchmark”?

- Benchmark is a program for performance measurement

## Requirements:

- Dimensions: throughput and latency
- Avoid significant overhead
- Test what is to be tested
- Perform a set of executions and provide stable reproducible results
- Should be easy to run



# Benchmark types

## ■ By scale

- Micro-benchmark (component level)
- Macro-benchmark (system level)

## ■ By nature

- Synthetic benchmark (emulate component load)
- Application benchmark (run real-world application)



# We'll talk about

- Synthetic micro-benchmark
  - Mimic component workload separately from the application
  - Measure performance of a small isolated piece of code
  
- The main concern
  - The smaller the Component we test – the stronger the impact of
    - Benchmark infrastructure overhead
    - JVM internal processes
    - OS and Hardware internals
    - ... and the phases of the Moon
  - Don't we really test one of those?



# When micro-benchmark is needed

- Most of the time it is not needed at all
- Does algorithm A work faster than B?  
(Consider equal analytical estimation)
- Does this tiny modification make any difference?  
(from Java, JVM, native code or hardware point of view)

# Tools needed to measure performance







# You had one job...

```
final int COUNT = 100;

long start = System.currentTimeMillis();
for (int i = 0; i < COUNT; i++) {
    // doStuff();
}
long duration = System.currentTimeMillis() - start;

long avg = duration / COUNT;
System.out.println("Average execution time is " + avg + " ms");
```



# Pitfall #0

- Using profiler to measure performance of small methods (adds significant overhead, measures execution “as is”)
- “You had one job” approach is enough in real life (not for micro-benchmarks, we got it already)
- Annotations and reflective benchmark invocations (you must be great at *java.lang.reflect* measurement)



# Micro-benchmark frameworks

- **JMH** – Takes into account a lot of internal VM processes and executes benchmarks with minimal infrastructure (Oracle)
- **Caliper** – Allows to measure repetitive code, works for Android, allows to post results online (Google)
- **Japex** – Allows to reduce infrastructure code, generates nice HTML reports with JFreeChart plots
- **JUnitPerf** – Measure functionality of the existing JUnit tests



# Java time interval measurement

## ■ *System.currentTimeMillis()*

- Value in milliseconds, but granularity depends on the OS
- Represents a “wall-clock” time (since the start of Epoch)

## ■ *System.nanoTime()*

- Value in nanoseconds, since some time offset
- The accuracy is not worse than *System.currentTimeMillis()*

## ■ *ThreadMXBean.getCurrentThreadCpuTime()*

- The actual CPU time spent for the thread (nanoseconds)
- Might be unsupported by your VM
- Might be expensive
- Relevant for a single thread only



Code warm-up, what happens before  
the steady-state





# Code warm-up, class loading

- A single warm-up iteration is NOT enough for Class Loading (not all the branches of classes may load on the first iteration)
- Sometimes classes are unloaded (it would be a shame if something messed your results up with a huge peak)
- Get help between iterations from
  - *ClassLoaderMXBean.getTotalLoadedClassCount()*
  - *ClassLoaderMXBean.getUnloadedClassCount()*
  - `-verbose:class`



# Code warm-up, compilation

- Classes are loaded, verified ~~and then being compiled~~
- *Oracle HotSpot* and *Azul Zing* run application in **interpreter**
- The hot method is being compiled after  
~10k (server), ~1.5k (client) invocations
- Long methods with loops are likely to be compiled earlier
  
- Check *CompilationMXBean.getTotalCompilationTime*
- Enable compilation logging with
  - -XX:+UnlockDiagnosticVMOptions
  - -XX:+PrintCompilation
  - -XX:+LogCompilation -XX:LogFile=<filename>



# Code warm-up, OSR

- Normal compilation and OSR will result in a similar code
- ...unless compiler is not able to optimize a given frame (e.g. inner loop is compiled before the outer one)
  
- In the real world normal compilation is more likely to happen, so it's better to avoid OSR in your benchmark
  - Do a set of small warm-up iterations instead of a single big one
  - Do not perform warm-up loops in the steady-state testing method



# Code warm-up, OSR example

Before:

```
public static void main(String... args) {  
    loop1: if(P1) goto done1  
        i=0;  
        loop2: if(P2) goto done2  
            A[i++];  
        goto loop2; // OSR goes here  
    done2:  
    goto loop1;  
    done1:  
}
```

After:

```
void OSR_main() {  
    A= // from interpreter  
    i= // from interpreter  
    loop2: if(P2) {  
        if(P1) goto done1  
        i=0;  
    } else { A[i++]; }  
    goto loop2  
    done1:  
}
```

- Now forget about array range check elimination




# Reaching the steady-state, summary

- Always do warm-up to reach steady-state
  - Use the same data and the same code
  - Discard warm-up results
  - Avoid OSR
  - Don't run benchmark in the “mixed modes” (interpreter/compiler)
  - Check class loading and compilation



# Using JMH

- Provides Maven archetype for a quick project setup
- Annotate your methods with `@GenerateMicroBenchmark`
- `mvn install` will build ready to use runnable jar with your benchmarks and needed infrastructure
- `java -jar target/mb.jar <benchmark regex> [options]`
- Will perform warm-up following by a set of iterations
- Print the results



Side effects that can affect performance





# Synchronization puzzle

```
void testSynchInner() {  
    synchronized (this) {  
        i++;  
    }  
}
```

8,244,087 usec

```
synchronized void testSynchOuter() {  
    i++;  
}
```

13,383,707 usec

# Synchronization puzzle, side effect

- Biased Locking: an optimization in the VM that leaves an object as logically locked by a given thread even after the thread has released the lock (cheap reacquisition)
- Does not work on VM start up  
(4 sec in HotSpot)  
Use `-XX:BiasedLockingStartupDelay=0`



# JVM optimizations (Good or Evil?)

- **WARNING:** some of the following optimizations will not work (at least for the given examples) in Java 6 (jdk1.6.0\_26), consider using Java 7 (jdk1.7.0\_21)





# Dead code elimination

- VM optimization eliminates dead branches of code
- Even if the code is meant to be executed, but the result is never used and does not have any side effect
- Always consume all the results of your benchmarked code
- Or you'll get the “*over 9000*” performance level
- Do not accumulate results or store them in class fields that are never used either
- Use them in the unobvious logical expression instead



# Dead code elimination, example

- Measurement: average nanoseconds / operation, less is better

```
private double n = 10;

1.008 public void stub() { }

1.017 public void dead() {
    @SuppressWarnings("unused")
    double r = n * Math.Log(n) / 2;
}

48.514 public void alive() {
    double r = n * Math.Log(n) / 2;
    if(r == n && r == 0)
        throw new IllegalStateException();
}
```



# Constant folding

- If the compiler sees that the result of calculation will always be the same, it will be stored in the constant value and reused
- Measurement: average nanoseconds / operation, less is better

	<code>private double x = Math.PI;</code>
1.014	<code>public void stub() { }</code>
	<code>public double wrong() {</code>
1.695	<code>    return Math.Log(Math.PI);</code>
	<code>}</code>
	<code>public double measureRight() {</code>
43.435	<code>    return Math.Log(x);</code>
	<code>}</code>



# Loop unrolling

- Is there anything bad?
- Measurement: average nanoseconds / operation, less is better

```
private double[] A = new double[2048];
```

```
public double plain() {  
    double sum = 0;  
    for (int i = 0; i < A.length; i++)  
        sum += A[i];  
    return sum;  
}
```

2773.883

```
public double manualUnroll() {  
    double sum = 0;  
    for (int i = 0; i < A.length; i += 4)  
        sum += A[i] + A[i + 1] + A[i + 2] + A[i + 3];  
    return sum;  
}
```

816.791



# Loop unrolling and hoisting

- Something bad happens when the loops of benchmark infrastructure code are unrolled
- And the calculations that we try to measure are hoisted from the loop

- For example, **Caliper** style benchmark looks like

```
private int reps(int reps) {  
    int s = 0;  
    for (int i = 0; i < reps; i++)  
        s += (x + y);  
    return s;  
}
```



# Loop unrolling and hoisting, example

```
@GenerateMicroBenchmark
public int measureRight() {
    return (x + y);
}
```

```
@GenerateMicroBenchmark
@OperationsPerInvocation(1)
public int measureWrong_1() {
    return reps(1);
}
```

...

```
@GenerateMicroBenchmark
@OperationsPerInvocation(N)
public int measureWrong_N() {
    return reps(N);
}
```



# Loop unrolling and hoisting, example

- Measurement: average nanoseconds / operation, less is better

Method	Result
Right	2.104
Wrong_1	2.055
Wrong_10	0.267
Wrong_100	0.033
Wrong_1000	0.057
Wrong_10000	0.045
Wrong_100000	0.043

# A word about concurrency

- Processes and threads fight for resources (single threaded benchmark is a utopia)





# Concurrency problems of benchmarks

- Benchmark states should be correctly
  - Initialized
  - Published
  - Shared between certain group of threads
- Multi threaded benchmark iteration should be synchronized and all threads should start their work at the same time
- No need to implement this infrastructure yourself, just write a correct benchmark using your favorite framework



# Full example, “human factor” included





# List iteration

- Which list implementation is faster for the foreach loop?
- ArrayList and LinkedList sequential iteration is linear,  $O(n)$ 
  - *ArrayList Iterator.next()*: return array[cursor++];
  - *LinkedList Iterator.next()*: return current = current.next;
- Let's check for the list of 1 million Integer's



# List iteration, foreach vs iterator

- Measurement: average milliseconds / operation, less is better

23.659

```
public List<Integer> arrayListForeach() {  
    for(Integer i : arrayList) {  
    }  
    return arrayList;  
}
```

22.445

```
public Iterator<Integer> arrayListIterator() {  
    Iterator<Integer> iterator = arrayList.iterator();  
    while(iterator.hasNext()) {  
        iterator.next();  
    }  
    return iterator;  
}
```



# List iteration, foreach < iterator, why?

- Foreach variant assigns element to a local variable  
`for(Integer i : arrayList)`
- Iterator variant does not  
`iterator.next();`
- We need to change Iterator variant to  
`Integer i = iterator.next();`
- So now it's correct to compare the results, at least according to the bytecode 😊



# List iteration, benchmark

```
@GenerateMicroBenchmark(BenchmarkType.ALL)  
public List<Integer> arrayListForeach() {  
    for(Integer i : arrayList) {  
    }  
    return arrayList;  
}
```

```
@GenerateMicroBenchmark(BenchmarkType.ALL)  
public Iterator<Integer> arrayListIterator() {  
    Iterator<Integer> iterator = arrayList.iterator();  
    while(iterator.hasNext()) {  
        Integer i = iterator.next();  
    }  
    return iterator;  
}
```

# List iteration, benchmark, result

- Measurement: average milliseconds / operation, less is better

List impl	Iteration	Java 6	Java 7
ArrayList	foreach	24.792	5.118
	iterator	24.769	0.140
LinkedList	foreach	15.236	9.485
	iterator	15.255	9.306

- Java 6 ArrayList uses AbstractList.Itr, LinkedList has its own, so there is less abstractions (in Java 7 ArrayList has its own optimized iterator)

# List iteration, benchmark, result

- Measurement: average milliseconds / operation, less is better

List impl	Iteration	Java 6	Java 7
ArrayList	foreach	24.792	5.118
	iterator	24.769	<b><u>0.140</u> WTF?!</b>
LinkedList	foreach	15.236	9.485
	iterator	15.255	9.306

# List iteration, benchmark, loop-hoisting

```
ListBenchmark.arrayListIterator()
```

```
Iterator<Integer> iterator = arrayList.iterator();  
while(iterator.hasNext()) {  
    iterator.next();  
}  
return iterator;
```

```
ArrayList.Itr<E>.next()
```

```
if (modCount != expectedModCount) throw new CME();  
int i = cursor;  
if (i >= size) throw new NoSuchElementException();  
Object[] elementData = ArrayList.this.elementData;  
if (i >= elementData.length) throw new CME();  
cursor = i + 1;  
return (E) elementData[lastRet = i];
```





# List iteration, benchmark, BlackHole

```
@GenerateMicroBenchmark(BenchmarkType.ALL)  
public void arrayListForeach(BlackHole bh) {  
    for(Integer i : arrayList) {  
        bh.consume(i);  
    }  
}
```

```
@GenerateMicroBenchmark(BenchmarkType.ALL)  
public void arrayListIterator(BlackHole bh) {  
    Iterator<Integer> iterator = arrayList.iterator();  
    while(iterator.hasNext()) {  
        Integer i = iterator.next();  
        bh.consume(i);  
    }  
}
```

# List iteration, benchmark, correct result

■ Measurement: average milliseconds / operation, less is better

List impl	Iteration	Java 6	Java 7	Java 7 BlackHole
ArrayList	foreach	24.792	5.118	8.550
	iterator	24.769	0.140	8.608
LinkedList	foreach	15.236	9.485	11.739
	iterator	15.255	9.306	11.763



# A fleeting glimpse on the JMH details

- We already know that JMH
  - Uses maven
  - Uses annotation-driven approach to detect benchmarks
  - Provides BlackHole to consume results (and CPU cycles)



# JMH: Building infrastructure

- Finds annotated micro-benchmarks using reflection
- Generates infrastructure plain java source code around the calls to the micro-benchmarks
- Compile, pack, run, profit
- No reflection during benchmark execution



# JMH: Various metrics

- Single execution time
- Operations per time unit
- Average time per operation
- Percentile estimation of time per operation



# JMH: Concurrency infrastructure

- @State of benchmark data is shared across the benchmark, thread, or a group of threads
- Allows to perform Fixtures (setUp and tearDown) in scope of the whole run, iteration or single execution
- @Threads a simple way to run concurrent test if you defined correct @State
- @Group threads to assign them for a particular role in the benchmark



# JMH: VM forking

- Allows to compare results obtained from various instances of VM
  - First test will work on the clean JVM and others will not
  - VM processes are not determined and may vary from run to run (compilation order, multi-threading, randomization)



# JMH: @CompilerControl

- Instructions whether to compile method or not
- Instructions whether to inline methods
- Inserting breakpoints into generated code
- Printing methods assembly






# Conclusions

- Do not reinvent the wheel, if you are not sure how it should work (consider using existing one)
- Consider the results being wrong if you don't have a clear explanation. Do not swallow that mystical behavior



Thanks for you attention



Questions?